

Introduction to Relational Databases

Principles of Data Science with R

Dr. Uma Ravat

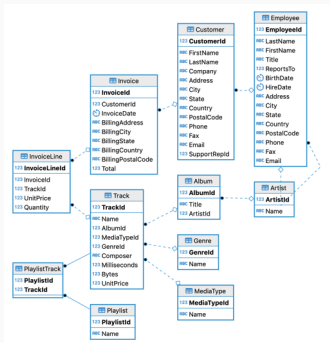
PSTAT 10

Next we will see. . .

- Intro to Databases
- Connecting to database from R
 - the R packages RSQLite, DBI
 - the database on disk/file `Chinook_Sqlite.sqlite`
- SQL (Structured Query Language)
- 3 parts of the relational data model

What is a (relational) database?

- A **database** is a **collection of tables** that are **related** to one another, together **with specification of these relations**.



- An **entity relationship diagram** (ER diagram) helps visualize the structure of the tables and their relation to one another in a database.

dataframes in R to Tables in databases

R jargon	Database jargon	definition
column	field/attribute	a variable/quantity of interest
row	record/tuple	collection of fields/attributes
dataframe	table/relation	a collection of records which all have the same fields/attribute (with different values)
types of the columns	table schema	datatype and other specifics about each field/attribute.

Students
Table
(data.frame)

Field
(Column)

id	firstname	lastname	age	instate
1	Billy	Joe	23	FALSE
2	Theodore	Squirrel	25	TRUE
3	Keeya	Nod	21	TRUE

Record
(Row)

Integer	Character	Character	Integer	Logical
---------	-----------	-----------	---------	---------

Table Scheme
(Data Types)

Why do we need database software?

- **Size**

- R keeps its data frames in memory
- Industrial databases are much bigger, need to store out of memory and bring into memory only required subsets
- Must work with selected subsets

- **Speed**

- Relational database model published in 70's by E. F. Codd at IBM labs in San Jose, CA.
- Smart people have worked very hard making relational databases efficient
- 2014 Turing award winner Michael Stonebraker

- **Concurrency**

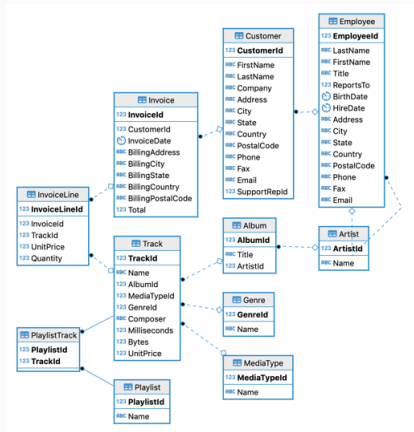
- Many users access the same database simultaneously
- Potential for trouble (two users want to change the same record at once)
- Database software takes care of this issue.

A **Relational Database Management System (RDBMS)** is software that enables you to create and maintain databases and interact with the data stored in tables in the relational database.

- **SQLite** is a popular light-weight, fast, full-featured RDBMS designed for simple applications (mobile apps)
 - is a simpler, file-based system that we will use via RSQLite package in R

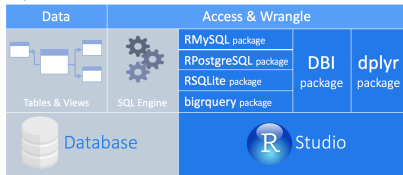
Chinook DB

This DB represents a digital media store, including tables for artists, albums, media tracks, invoices, and customers.



Connecting R to SQLite

Open Source Databases



- First, we need to install the packages DBI, RSQLite,

```
install.packages("DBI")
```

```
install.packages("RSQLite")
```

- then we load them into our R session with `library()`

Also, we need a SQLite database file available at `Chinook_Sqlite.sqlite` for this lecture

Connecting to the database in R

```
library(DBI)
library(RSQLite)
drv = dbDriver("SQLite")
chinook_db = dbConnect(drv,
                        dbname="./data/Chinook_Sqlite.sqlite")
```

The object `chinook_db` is now a persistent connection to the `Chinook_Sqlite.sqlite` database on disk.

Listing what's available

The data in a relational database is stored in relations, aka **tables**:

```
# List tables in our database
```

```
dbListTables(chinook_db)
```

```
## [1] "Album"      "Artist"     "Customer"   "Employee"  
## [5] "Genre"      "Invoice"    "InvoiceLine" "MediaType"  
## [9] "Playlist"   "PlaylistTrack" "Track"
```

Each table has rows of tuples, aka **records**, and columns of attributes, aka **fields**.

```
#List fields in Customer table
```

```
dbListFields(chinook_db, "Customer")
```

```
## [1] "CustomerId" "FirstName" "LastName" "Company" "Address"
```

```
## [6] "City" "State" "Country" "PostalCode" "Phone"
```

```
## [11] "Fax" "Email" "SupportRepId"
```

Importing a table as a data frame

```
customer = dbReadTable(chinook_db, "Customer")  
class(customer)
```

```
## [1] "data.frame"
```

```
dim(customer)
```

```
## [1] 60 13
```

Now we could go on and perform R operations on `customer`, since it's a dataframe

- **SQL (structured query language)** is language/software to interact with data stored in databases
- **SQL** is the standard for database software
 - many different implementations of SQL, each with unique features.
- Most basic actions with databases are **SQL queries**, like row/column selections, inserts, updates and deletes

SELECT

Main tool in the SQL language: **SELECT**, which allows you to perform queries on a particular table in a database. It has the form:

```
SELECT columns  
  FROM table  
  WHERE condition  
  GROUP BY columns  
  HAVING condition  
  ORDER BY column [ASC | DESC]  
  LIMIT offset, count;
```

WHERE, GROUP BY, HAVING, ORDER BY, LIMIT are all optional

Example

Pick out five columns from the table “Customer”, and only look at the first 6 rows:

```
dbGetQuery(chinook_db,  
            "select CustomerId, FirstName, LastName, City, Country  
            from Customer  
            limit 6" )
```

##	CustomerId	FirstName	LastName	City	Country
## 1	1	Luís	Gonçalves	São José dos Campos	Brazil
## 2	2	Leonie	Köhler	Stuttgart	Germany
## 3	3	François	Tremblay	Montréal	Canada
## 4	4	Bjørn	Hansen	Oslo	Norway
## 5	5	František	Wichterlová	Prague	Czech Republic
## 6	6	Helena	Holý	Prague	Czech Republic

To replicate this simple command on the imported data frame:

```
customer[1:6, c("CustomerId", "FirstName", "LastName", "City", "Country")]
```

##	CustomerId	FirstName	LastName	City	Country
## 1	1	Luís	Gonçalves	São José dos Campos	Brazil
## 2	2	Leonie	Köhler	Stuttgart	Germany
## 3	3	François	Tremblay	Montréal	Canada
## 4	4	Bjørn	Hansen	Oslo	Norway
## 5	5	František	Wichterlová	Prague	Czech Republic
## 6	6	Helena	Holý	Prague	Czech Republic

ORDER BY

We can use the ORDER BY option in SELECT to specify an ordering for the rows

Default is ascending order; add DESC for descending

```
dbGetQuery(chinook_db,  
            "select CustomerId, FirstName, LastName, City, Country  
              from Customer  
              ORDER BY FirstName DESC  
              limit 10")
```

##	CustomerId	FirstName	LastName	City	Country
## 1	42	Wyatt	Girard	Bordeaux	France
## 2	25	Victor	Stevens	Madison	USA
## 3	19	Tim	Goyer	Cupertino	USA
## 4	44	Terhi	Hämäläinen	Helsinki	Finland
## 5	54	Steve	Murray	Edinburgh	United Kingdom
## 6	49	Stanisław	Wójcik	Warsaw	Poland
## 7	12	Roberto	Almeida	Rio de Janeiro	Brazil
## 8	29	Robert	Brown	Toronto	Canada
## 9	26	Richard	Cunningham	Fort Worth	USA
## 10	59	Puja	Srivastava	Bangalore	India

Close connection

We opened a connection as follows:

```
chinook_db <- dbConnect(SQLite(),  
                        "Chinook_Sqlite.sqlite")
```

After the end of a session, it is good practice to explicitly close your connection.

```
dbDisconnect(chinook_db)
```

If indeed the connection is closed, reading some data should give an error

```
dbGetQuery(chinook_db, "select CustomerId, FirstName, LastName from Customer")
```

Error in h(simpleError(msg, call)) : error in evaluating the argument 'conn' in selecting a method for function 'dbGetQuery': object 'chinook_db' not found

Back to some theory: The relational model

In the **relational database**, the data is organized into **relations/tables**.

A relational model is used to represent data and the relationships between data items



The Relational Model consist of 3 parts

1. Manipulative part

- Allows for data to be manipulated in the database.
- SQL for create, update, delete tables, databases and user access
- SQL for select, insert, update, delete data in tables
- We saw some SQL queries for Select; more later

2. Structural part

- Tables, relations between tables and rules/constraints for these
- Visually as ER diagram: Table schema and relations between tables
- Database Schema and rules/constraints, Primary Keys and Foreign Keys.

3. Integrity part

- Rules to maintain integrity of data
- Necessary to keep data complete, consistent and reliable.
 - Entity integrity: integrity of each relation
 - Referential integrity : integrity between relations

Table/Relation

- Each relation must have a **unique name** in the database
- **No duplicate rows**, or tuples, are allowed in a table/relation
- **Each row** in a table/relation has **its own unique key**

Why Keys?

Students					Grades		
id	firstname	lastname	age	instate	student_id	course_id	grade
1	Bitly	Jones	19	FALSE	1	1	A
2	Tiny	Lark	25	TRUE	2	2	A
3	Hugh	Grand	22	TRUE	1	2	B
					3	4	A
					3	1	B

Courses			Exams		
course_id	Name	Department	student_id	name	exam
1	PSTAT10	pstat	1	PSTAT10	90.9
2	PSTAT120A	pstat	1	PSTAT120A	84.5
3	PSTAT120B	pstat	2	PTAT120A	90.7
4	HIST101	hist	3	HIST101	83
5	HIST201	hist	3	PSTAT10	96

Why Keys?

	Grades				Exams	
student_id	course	grade		student_id	course	exam
1	PSTAT10	A		1	PSTAT10	90.9
2	PSTAT120A	A		1	PSTAT120A	84.5
1	PSTAT120A	B		2	PTAT120A	90.7
3	HIST101	A		3	HIST101	83
3	PSTAT10	B		3	PSTAT10	96

Why Keys?

	Grades				Exams	
student_id	course	grade		student_id	course	exam
1	PSTAT10	A		1	PSTAT10	90.9
2	PSTAT120A	A		1	PSTAT120A	84.5
1	PSTAT120A	B		2	PSTAT120A	90.7
3	HIST101	A		3	HIST101	83
3	PSTAT10	B		3	PSTAT10	96
	student_id	course	exam	grade		
	1	PSTAT10	90.9	A		
	2	PSTAT120A	84.5	A		
	1	PSTAT120A	90.7	B		
	3	HIST101	83	A		
	3	PSTAT10	96	B		

Structural part : Keys

A Key

- consists of one or more attributes
- places certain constraints on a databases
- Each row in a table/relation has its own unique key (primary key)
- used to establish and identify relationships between relations
 - a row in a relation can be linked with another row in other relations (foreign key)

4 types of keys

- Super key
- Candidate key
- Primary key
- Foreign key

Keys: Super Keys

A super key is a set of one or more attributes that uniquely identify a tuple in a relation.

APPOINTMENT	CustNmb	CustName	ApptDay
	5	Brian	Monday
	213	Grayson	Tuesday
	7	Jon	Monday
	88	Nitin	Wednesday
	7	Jon	Tuesday

Super Keys for the Appointment relation:

- {CustNmb, CustName, ApptDay},
- {CustNmb, ApptDay},
- {CustName, ApptDay}

What about {CustNmb, CustName}, {CustNmb}, {CustName}, {ApptDay} ?

Keys: Candidate Keys

- Candidate key is a minimal set of attributes which can uniquely identify a tuple.
- Candidate keys are selected from the set of super keys.
- Candidate keys should not have any redundant attributes.
- There can be more than one candidate key in a relation
- The candidate key can be simple, having only one attribute, or it can be a composite of multiple attributes.
- A candidate key is a super key, but not the other way round

Keys: Candidate Keys

APPOINTMENT	CustNmb	CustName	ApptDay
	5	Brian	Monday
	213	Grayson	Tuesday
	7	Jon	Monday
	88	Nitin	Wednesday
	7	Jon	Tuesday

Candidate Keys for the Appointment relation

- {CustNmb, ApptDay},
- {CustName, ApptDay}

What about {CustNmb, CustName, ApptDay}, {CustNmb, CustName}, {CustNmb}, {CustName}, {ApptDay} ?

Keys: Primary Keys

- A primary key is the minimal set of attributes which can uniquely identify a tuple.
- one of the candidate keys is chosen as the primary key.
- Candidate keys NOT selected for use as the primary key are called Alternate Keys.
- The same primary key cannot be used for different relations
- Primary key values cannot be null

Keys: Primary Keys

APPOINTMENT	CustNmb	CustName	ApptDay
	5	Brian	Monday
	213	Grayson	Tuesday
	7	Jon	Monday
	88	Nitin	Wednesday
	7	Jon	Tuesday

The primary key for our APPOINTMENT relation is

- {CustNmb, ApptDay}

What about {CustNmb, CustName, ApptDay}, {CustNmb, CustName}, {CustNmb}, {CustName}, {ApptDay} ?

What about {CustName, ApptDay}?

Back to Chinook DB : Field metadata

Unlike a dataframe, there is extra information in a database table that expresses relational information between tables.

```
dbGetQuery(chinook_db, "pragma table_info(Customer)")
```

##	cid	name	type	notnull	dflt_value	pk
## 1	0	CustomerId	INTEGER	1	NA	1
## 2	1	FirstName	NVARCHAR(40)	1	NA	0
## 3	2	LastName	NVARCHAR(20)	1	NA	0
## 4	3	Company	NVARCHAR(80)	0	NA	0
## 5	4	Address	NVARCHAR(70)	0	NA	0
## 6	5	City	NVARCHAR(40)	0	NA	0
## 7	6	State	NVARCHAR(40)	0	NA	0
## 8	7	Country	NVARCHAR(40)	0	NA	0
## 9	8	PostalCode	NVARCHAR(10)	0	NA	0
## 10	9	Phone	NVARCHAR(24)	0	NA	0
## 11	10	Fax	NVARCHAR(24)	0	NA	0
## 12	11	Email	NVARCHAR(60)	1	NA	0
## 13	12	SupportRepId	INTEGER	0	NA	0

KEYS: Primary key

The **primary key** is a *unique identifier* of the rows in a table. Two rows cannot have the same primary key:

```
dbGetQuery(chinook_db,  
  "select CustomerId, FirstName, LastName, City, Country  
    from Customer  
   limit 2")
```

##	CustomerId	FirstName	LastName	City	Country
## 1	1	Luís	Gonçalves	São José dos Campos	Brazil
## 2	2	Leonie	Köhler	Stuttgart	Germany

```
dbExecute(chinook_db,  
  paste("insert into customer",  
    "(CustomerId, FirstName, LastName, Email)",  
    "values",  
    "(1, 'Luis','Armstrong','LuisArmstrong@pstat.ucsb.edu')"))
```

```
## Error: UNIQUE constraint failed: Customer.CustomerId
```

CustomerId is the **primary key** and must be unique.

Multi-column primary key

Primary key's can consist of multiple columns if it takes multiple columns to identify a row in a table. But, two rows cannot have the same primary key.

Single column primary key

```
dbGetQuery(chinook_db, "pragma table_info(Customer)")
```

##	cid	name	type	notnull	dflt_value	pk
## 1	0	CustomerId	INTEGER	1	NA	1
## 2	1	FirstName	NVARCHAR(40)	1	NA	0
## 3	2	LastName	NVARCHAR(20)	1	NA	0
## 4	3	Company	NVARCHAR(80)	0	NA	0
## 5	4	Address	NVARCHAR(70)	0	NA	0
## 6	5	City	NVARCHAR(40)	0	NA	0
## 7	6	State	NVARCHAR(40)	0	NA	0
## 8	7	Country	NVARCHAR(40)	0	NA	0
## 9	8	PostalCode	NVARCHAR(10)	0	NA	0
## 10	9	Phone	NVARCHAR(24)	0	NA	0
## 11	10	Fax	NVARCHAR(24)	0	NA	0
## 12	11	Email	NVARCHAR(60)	1	NA	0
## 13	12	SupportRepId	INTEGER	0	NA	0

Multi column primary key

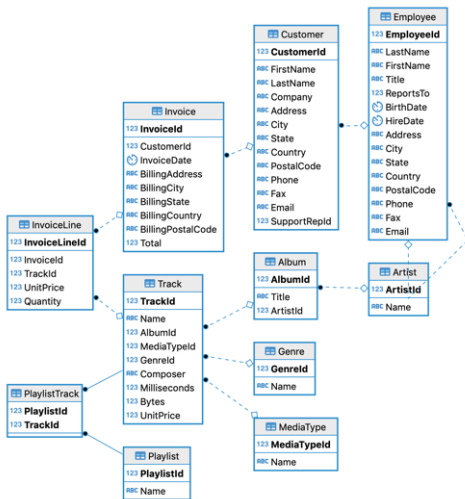
```
dbGetQuery(chinook_db, "pragma table_info(PlaylistTrack)")
```

##	cid	name	type	notnull	dflt_value	pk
## 1	0	PlaylistId	INTEGER	1	NA	1
## 2	1	TrackId	INTEGER	1	NA	2

KEYS: Primary key

Tables are not required to have a primary key, but most do. All the tables in Chinook have a primary key.

In ER diagram, Primary Key is denoted as **bold** or underlined field name.



KEYS: Foreign Keys

The foreign key is a set of attributes that matches the primary key for another relation. A foreign key field *points to* the primary key of another table.

The foreign key links two relations and can be used to cross-reference the relations.



KEYS: Foreign Keys

The relationship between tables is expressed by primary keys and **foreign keys**.

Remember we are working with a relational database, following a relational data model.

```
dbGetQuery(chinook_db,  
            "pragma foreign_key_list(customer)")
```

```
##   id seq    table          from          to on_update on_delete match  
## 1  0    0 Employee SupportRepId EmployeeId NO ACTION NO ACTION  NONE
```

KEYS: Enforcing Foreign keys

Foreign keys must either point to an existing value or be NULL.

To enforce Foreign key constraints in SQLite RDBMS

Required for foreign-key support otherwise foreign keys are not enforced

```
dbExecute(chinook_db, "pragma foreign_keys = on")
```

```
dbGetQuery(chinook_db,  
            "SELECT max(EmployeeId) FROM Employee")
```

```
##      max(EmployeeId)  
## 1                8
```

```
dbGetQuery(chinook_db,  
            "SELECT max(CustomerId) FROM Customer")
```

```
##      max(CustomerId)  
## 1             888
```

```
dbExecute(chinook_db,  
  "INSERT INTO Customer  
    (CustomerId, FirstName, LastName, Email, SupportRepId)  
  VALUES  
    (59, 'Luis', 'Armstrong', 'luisArmstrong@pstat.ucsb.edu', 88)")
```

```
## Error: UNIQUE constraint failed: Customer.CustomerId
```

```
dbExecute(chinook_db,  
  "INSERT INTO Customer  
    (CustomerId, FirstName, LastName, Email, SupportRepId)  
  VALUES  
    (60, 'Luis', 'Armstrong', 'luisArmstrong@pstat.ucsb.edu', 10)")
```

```
## Error: FOREIGN KEY constraint failed
```


Interpretation of foreign key



- Each customer in Customer table *can* be assigned a support representative
- The support rep is an employee at the store and therefore has a unique id, EmployeeId
- This unique id, EmployeeId, is the primary key of the employee table

Thus real-world relationship is encoded by the relational model using primary and foreign key relationships.

Keys: Why do we need keys?

- For identifying any row of data in a relation uniquely
- They ensure integrity of data is maintained.
- Keys establish relationships between relations and identify relationships between relations

The Relational Model consist of 3 parts

1. Manipulative part

- Allows for data is manipulated in the database.
- SQL for create, update, delete tables, databases and user access
- SQL for select, insert, update, delete data in tables
- We saw some SQL queries for Select; more later

2. Structural part

- Tables, relations between tables and rules/constraints for these
- Visually as ER diagram: Table schema and relations between tables
- Database Schema and rules/constraints, Primary Keys and Foreign Keys.

3. Integrity part

- Rules to maintain integrity of data
- Necessary to keep data complete, consistent and reliable.
 - **Entity integrity:** integrity of each relation
 - **Referential integrity:** integrity between relations

Integrity Constraints

We have seen two examples of *integrity constraints*:

- **Entity integrity:** Primary keys must be unique (and not NULL)
 - This ensures there are no duplicate records
- **Referential integrity:** Foreign keys must reference existing primary keys or be NULL
 - ensures that cross-references to non-existing tuples cannot occur

These constraints enforce the *integrity* of a database; no bad data or corrupted relationships.

Keys help maintain the integrity of the data

Database Schema

The **schema** of a database describes its *structure*:

- Names of all the tables
- Names of all fields in each table
- Primary key/foreign key relationships between tables
- Other metadata (data types of each field in each table, ...)

Basically everything other than the actual data itself.

Represented via E-R diagrams (Entity relationship)

We have been looking at parts of the schema with the `pragma` keyword.

```
dbGetQuery(chinook_db, "pragma table_info(customer)")
```

##	cid	name	type	notnull	dflt_value	pk
## 1	0	CustomerId	INTEGER	1	NA	1
## 2	1	FirstName	NVARCHAR(40)	1	NA	0
## 3	2	LastName	NVARCHAR(20)	1	NA	0
## 4	3	Company	NVARCHAR(80)	0	NA	0
## 5	4	Address	NVARCHAR(70)	0	NA	0

- Databases are used to store massive amounts of data that cannot fit in memory.
- SQL is the language used to manipulate relational databases
- SQLite is the SQL implementation we will use, provided by the RSQLite package.